

Part 1:

```
# Import necessary libraries
import torch                # PyTorch library for building and
training neural networks
import torch.nn as nn       # Neural network module
import torch.optim as optim  # Optimization algorithms
import numpy as np          # Numerical operations
import matplotlib.pyplot as plt # Visualization
from torch.utils.data import DataLoader, TensorDataset # Data
handling utilities

# ----- Dataset Preparation -----

# Generate a circular dataset (synthetic data points on a circle)
theta = np.linspace(0, 2 * np.pi, 100) # Angles from 0 to 2π
radius = 3                               # Fixed radius for the circle
x_circle = radius * np.cos(theta)        # x-coordinates of the circle
y_circle = radius * np.sin(theta)        # y-coordinates of the circle
data = np.stack((x_circle, y_circle), axis=1) # Combine x and y into
2D points

# Convert data to PyTorch tensors (for use in the model)
data = torch.tensor(data, dtype=torch.float32)

# Create a DataLoader to iterate over the data in batches
batch_size = 10 # Number of samples per batch
dataset_loader = DataLoader(TensorDataset(data),
batch_size=batch_size, shuffle=True)

# ----- VAE Model Definition -----

# Define the Variational Autoencoder (VAE) class
class VAE(nn.Module):
    def __init__(self, input_dim, latent_dim):
        super(VAE, self).__init__()

        # Encoder: Maps input data to a latent representation
        self.encoder = nn.Sequential(
            nn.Linear(input_dim, 32), # First layer with 32 neurons
            nn.ReLU(),                # Activation function
            nn.Linear(32, 16),         # Second layer with 16 neurons
            nn.ReLU()
        )
        self.fc_mu = nn.Linear(16, latent_dim) # Outputs the mean
        # (μ) of the latent space
        self.fc_logvar = nn.Linear(16, latent_dim) # Outputs the log-
        # variance (log σ^2) of the latent space
```

```

        # Decoder: Maps the latent representation back to the input
space
        self.decoder = nn.Sequential(
            nn.Linear(latent_dim, 16), # First layer in decoder
            nn.ReLU(),
            nn.Linear(16, 32),         # Second layer
            nn.ReLU(),
            nn.Linear(32, input_dim)   # Final layer to match input
dimensions
        )

        # Encode function: Projects input to latent space parameters ( $\mu$ ,
log  $\sigma^2$ )
        def encode(self, x):
            h = self.encoder(x)
            mu = self.fc_mu(h)          # Mean of the latent space
            logvar = self.fc_logvar(h)  # Log-variance of the latent
space
            return mu, logvar

        # Reparameterization trick: Samples latent variables  $z$  from  $N(\mu, \sigma^2)$ 
        def reparameterize(self, mu, logvar):
            std = torch.exp(0.5 * logvar) # Standard deviation from log-
variance
            eps = torch.randn_like(std)   # Random noise sampled from a
standard normal distribution
            return mu + eps * std         # Latent variable  $z = \mu + \sigma * \epsilon$ 

        # Decode function: Maps latent variable  $z$  back to input space
        def decode(self, z):
            return self.decoder(z)

        # Forward pass: Combines encoding, reparameterization, and
decoding
        def forward(self, x):
            mu, logvar = self.encode(x)    # Encode input
            z = self.reparameterize(mu, logvar) # Sample  $z$  using
reparameterization trick
            return self.decode(z), mu, logvar # Return reconstructed
input,  $\mu$ , and log  $\sigma^2$ 

# ----- Loss Function -----

# Define the VAE loss function
def vae_loss(recon_x, x, mu, logvar, recon_weight=10):
    # Reconstruction loss: Measures how well the reconstructed data
matches the original

```

```

    recon_loss = recon_weight * nn.MSELoss()(recon_x, x)

    # KL divergence: Regularizes the latent space to follow a standard
    normal distribution
    kl_div = -0.05 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())

    return recon_loss + kl_div / x.size(0) # Normalize KL divergence
    by batch size

# ----- Model Initialization -----

# Define model parameters
input_dim = 2 # Dimensionality of the input (x, y coordinates)
latent_dim = 2 # Latent space dimension (allows 2D representation of
the data)

# Initialize the VAE model and optimizer
vae = VAE(input_dim, latent_dim)
optimizer = optim.Adam(vae.parameters(), lr=0.001) # Adam optimizer
with learning rate 0.001

# ----- Training the VAE -----

# Define training parameters
num_epochs = 1000 # Number of epochs for training
for epoch in range(num_epochs):
    vae.train() # Set model to training mode
    total_loss = 0

    for batch in dataset_loader:
        batch_data = batch[0] # Extract batch data
        optimizer.zero_grad() # Clear gradients from the
        previous step

        # Forward pass through the VAE
        recon_data, mu, logvar = vae(batch_data)

        # Compute the VAE loss
        loss = vae_loss(recon_data, batch_data, mu, logvar)

        # Backpropagate and update model parameters
        loss.backward()
        optimizer.step()

        total_loss += loss.item() # Accumulate loss for the epoch

# Print average loss every 100 epochs
if (epoch + 1) % 100 == 0:
    avg_loss = total_loss / len(dataset_loader)

```

```

    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}')

Epoch [100/1000], Loss: 0.8002
Epoch [200/1000], Loss: 0.7296
Epoch [300/1000], Loss: 0.7088
Epoch [400/1000], Loss: 0.7117
Epoch [500/1000], Loss: 0.7242
Epoch [600/1000], Loss: 0.7210
Epoch [700/1000], Loss: 0.7103
Epoch [800/1000], Loss: 0.6901
Epoch [900/1000], Loss: 0.6882
Epoch [1000/1000], Loss: 0.7185

# ----- Evaluation and Visualization -----

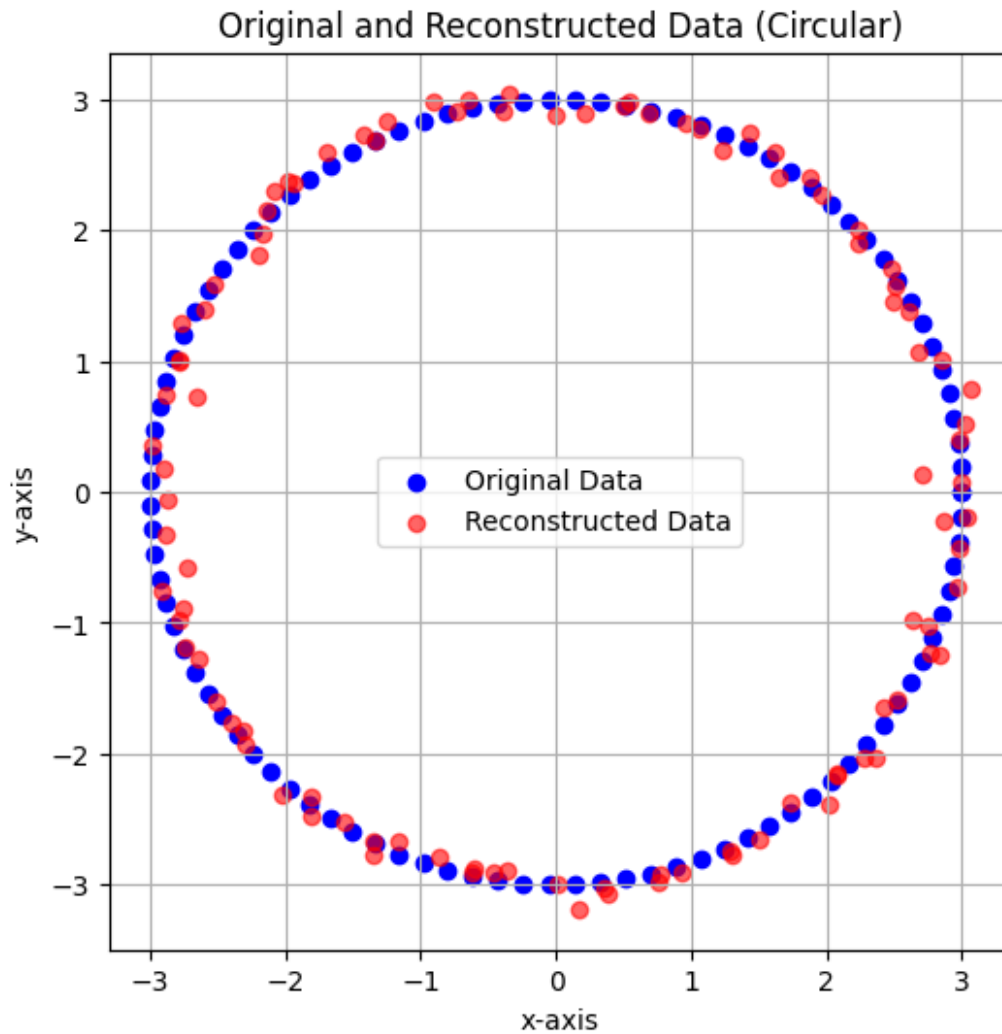
# Set the model to evaluation mode
vae.eval()
reconstructed_data = []

# Generate reconstructed data from the model
with torch.no_grad():
    for batch in dataset_loader:
        batch_data = batch[0]
        recon_batch, _, _ = vae(batch_data)
        reconstructed_data.append(recon_batch)

# Convert reconstructed data to numpy format for visualization
reconstructed_data = torch.cat(reconstructed_data).numpy()
data_np = data.numpy()

# Plot original data and reconstructed data
plt.figure(figsize=(6, 6))
plt.scatter(data_np[:, 0], data_np[:, 1], color='blue',
            label='Original Data') # Original circular data
plt.scatter(reconstructed_data[:, 0], reconstructed_data[:, 1],
            color='red', alpha=0.6, label='Reconstructed Data') # Reconstructed
data
plt.legend()
plt.title("Original and Reconstructed Data (Circular)")
plt.xlabel("x-axis")
plt.ylabel("y-axis")
plt.grid(True)
plt.show()

```



Part 2:

*# To plot the latent space of the trained VAE by extracting the `mu` (mean) values for each input point, you can modify the code as follows:*

*# ----- Latent Space Visualization -----*

*# Set the model to evaluation mode*

`vae.eval()`

`latent_space = []`

*# Extract the latent space (mu values) for each input point*

`with torch.no_grad():`

`for batch in dataset_loader:`

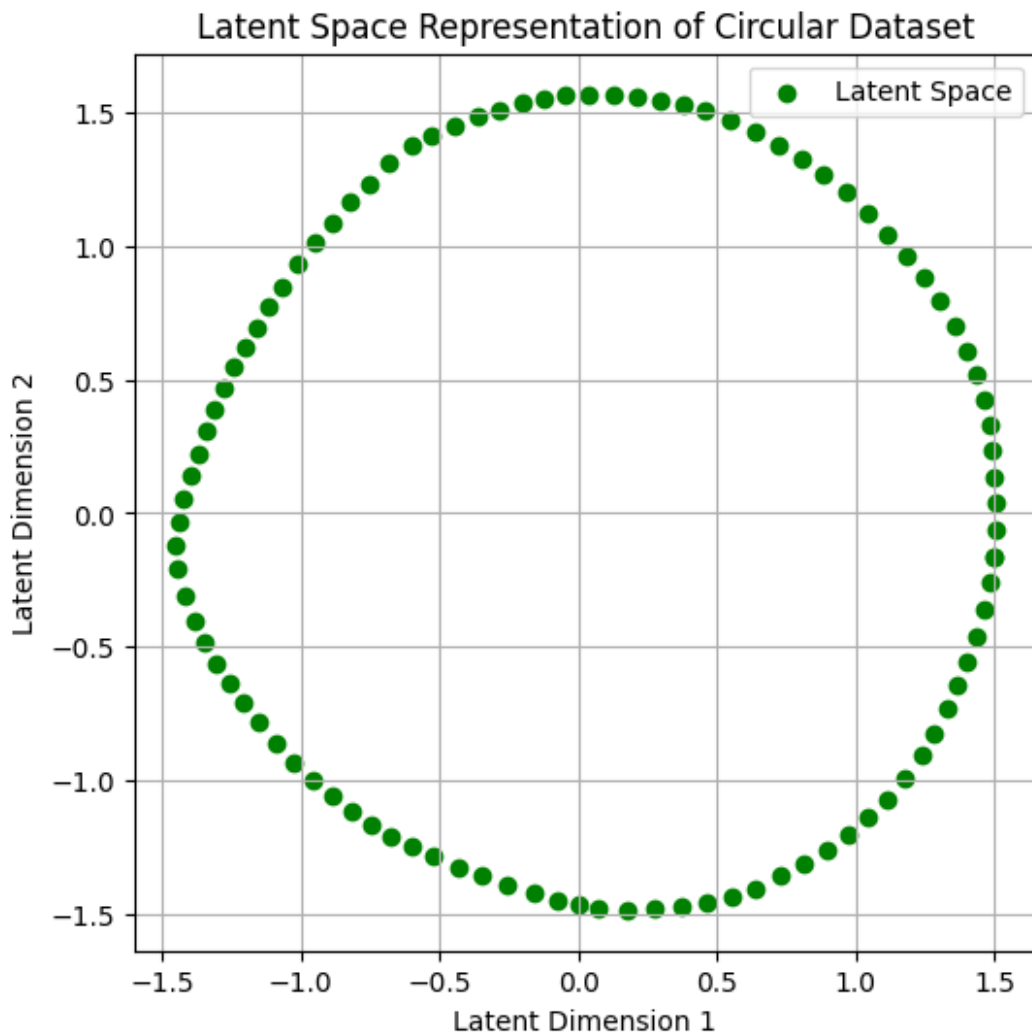
`batch_data = batch[0]`

`mu, _ = vae.encode(batch_data)`

`latent_space.append(mu)`

```
# Convert latent space data to numpy format for visualization
latent_space = torch.cat(latent_space).numpy()

# Plot the latent space
plt.figure(figsize=(6, 6))
plt.scatter(latent_space[:, 0], latent_space[:, 1], color='green',
            label='Latent Space')
plt.legend()
plt.title("Latent Space Representation of Circular Dataset")
plt.xlabel("Latent Dimension 1")
plt.ylabel("Latent Dimension 2")
plt.grid(True)
plt.show()
```



Part 3:

```

# ----- New Synthetic Data Preparation (Spiral)
# -----

# Generate a spiral dataset
theta_spiral = np.linspace(0, 4 * np.pi, 100) # Angles from 0 to 4π
r_spiral = theta_spiral # Radius increases with angle
x_spiral = r_spiral * np.cos(theta_spiral) # x-coordinates of the spiral
y_spiral = r_spiral * np.sin(theta_spiral) # y-coordinates of the spiral
data_spiral = np.stack((x_spiral, y_spiral), axis=1) # Combine x and y into 2D points

# Convert data to PyTorch tensors (for use in the model)
data_spiral = torch.tensor(data_spiral, dtype=torch.float32)

# Create a DataLoader to iterate over the data in batches
dataset_loader_spiral = DataLoader(TensorDataset(data_spiral),
batch_size=batch_size, shuffle=True)

# ----- Training the VAE on Spiral Data
# -----

# Define training parameters
num_epochs_spiral = 1000 # Number of epochs for training
for epoch in range(num_epochs_spiral):
    vae.train() # Set model to training mode
    total_loss_spiral = 0

    for batch in dataset_loader_spiral:
        batch_data = batch[0] # Extract batch data
        optimizer.zero_grad() # Clear gradients from the
previous step

        # Forward pass through the VAE
        recon_data, mu, logvar = vae(batch_data)

        # Compute the VAE loss
        loss = vae_loss(recon_data, batch_data, mu, logvar)

        # Backpropagate and update model parameters
        loss.backward()
        optimizer.step()

        total_loss_spiral += loss.item() # Accumulate loss for the
epoch

    # Print average loss every 100 epochs
    if (epoch + 1) % 100 == 0:
        avg_loss_spiral = total_loss_spiral /
len(dataset_loader_spiral)

```

```

        print(f'Epoch [{epoch+1}/{num_epochs_spiral}], Loss:
{avg_loss_spiral:.4f}')

# ----- Evaluation and Visualization of Spiral Data
-----

# Set the model to evaluation mode
vae.eval()
reconstructed_data_spiral = []

# Generate reconstructed data from the model
with torch.no_grad():
    for batch in dataset_loader_spiral:
        batch_data = batch[0]
        recon_batch, _, _ = vae(batch_data)
        reconstructed_data_spiral.append(recon_batch)

# Convert reconstructed data to numpy format for visualization
reconstructed_data_spiral =
torch.cat(reconstructed_data_spiral).numpy()
data_spiral_np = data_spiral.numpy()

# Plot original data and reconstructed data
plt.figure(figsize=(6, 6))
plt.scatter(data_spiral_np[:, 0], data_spiral_np[:, 1], color='blue',
label='Original Spiral Data') # Original spiral data
plt.scatter(reconstructed_data_spiral[:, 0],
reconstructed_data_spiral[:, 1], color='red', alpha=0.6,
label='Reconstructed Spiral Data') # Reconstructed data
plt.legend()
plt.title("Original and Reconstructed Data (Spiral)")
plt.xlabel("x-axis")
plt.ylabel("y-axis")
plt.grid(True)
plt.show()

Epoch [100/1000], Loss: 1.2961
Epoch [200/1000], Loss: 1.1519
Epoch [300/1000], Loss: 1.0709
Epoch [400/1000], Loss: 1.0176
Epoch [500/1000], Loss: 0.9981
Epoch [600/1000], Loss: 1.0884
Epoch [700/1000], Loss: 0.9332
Epoch [800/1000], Loss: 1.0384
Epoch [900/1000], Loss: 0.9453
Epoch [1000/1000], Loss: 0.9344

```



